

Towards Efficient Random-Order Enumeration for Join Queries

Technical Report

0.1 Proof of Lemma 2.

PROOF. We prove the contrapositive: if $\varphi_Q^*(i) \neq \text{false}$, then $i \leq \text{upp}(\psi_r)$. By the definition of φ^* , for each integer $i > 0$ satisfying $\varphi_Q^*(i) \neq \text{false}$, there is a tuple $t \in \text{Res}(Q)$ along with a root-to-leaf path u_1, \dots, u_h ($u_1 = r, u_h = u_t$) such that (11) holds. Then, let

$$\begin{aligned} \text{sum}_j &= \sum_{\substack{\psi \prec \psi_{u_{j+1}} \\ \psi \in \text{children}(\psi_{u_j})}} \text{upp}(\psi) \\ &\leq \sum_{\psi \in \text{children}(\psi_{u_j})} \text{upp}(\psi) - \text{upp}(\psi_{u_{j+1}}) \\ &\leq \text{upp}(\psi_{u_j}) - \text{upp}(\psi_{u_{j+1}}), \end{aligned} \quad (1)$$

therefore

$$\begin{aligned} i &= \sum_{j=1}^{h-1} \text{sum}_j + 1 \leq \sum_{j=1}^{h-1} (\text{upp}(\psi_{u_j}) - \text{upp}(\psi_{u_{j+1}})) + 1 \\ &= \text{upp}(\psi_{u_1}) - \text{upp}(\psi_{u_h}) + 1 = \text{upp}(\psi_r). \end{aligned} \quad (2)$$

□

0.2 Proof of Lemma 4.

PROOF. For each tuple $t \in R$ satisfying ψ , we have $l_i \leq t(x_i) \leq h_i$ for $\forall i \in \{r_1, \dots, r_d\}$. This implies $t_l^\psi \preceq t \preceq t_h^\psi$, and consequently $t \in R[t_l^\psi, t_h^\psi]$. Therefore, $R|_\psi \subseteq R[t_l^\psi, t_h^\psi]$. For each tuple $t \in R[t_l^\psi, t_h^\psi]$, we have $t_l^\psi \preceq t \preceq t_h^\psi$. We need to prove that for each $x_i \in \text{att}(R)$, $l_i \leq t(x_i) \leq h_i$. Let s be a split position of ψ , then:

- (1) For $\forall x_i \in \text{att}(R)$ with $\forall 1 \leq i < s$, since $l_i = h_i$, it follows that $l_i = t(x_i) = h_i$. Otherwise, one of the conditions $t \succeq t_l^\psi$ or $t \preceq t_h^\psi$ would be violated.
- (2) If $x_s \in \text{att}(R)$ then $l_s \leq t(x_s) \leq h_s$, otherwise $t_l^\psi \preceq t \preceq t_h^\psi$ would be violated.
- (3) For $\forall x_i \in \text{att}(R)$ with $\forall s < i \leq n$, since $l_i = \min_Q(x_i)$ and $h_i = \max_Q(x_i)$, it follows that $l_i \leq t(x_i) \leq h_i$.

Then t satisfies ψ , i.e., $t \in R|_\psi$. Therefore, $R[t_l^\psi, t_h^\psi] \subseteq R|_\psi$ and then we have $R[t_l^\psi, t_h^\psi] = R|_\psi$, which implies that $R.\text{cnt}(\psi) = |R|_\psi$. □

0.3 Proof of Lemma 6.

PROOF. Since $\text{upp}(\psi) \leq 1$, two cases need to be discussed. Firstly, for the case $\text{upp}(\psi) = 0$, we have $|\text{Res}(Q|_\psi)| \leq \text{upp}(\psi) = 0$, which implies that $\text{Res}(Q|_\psi) = \emptyset$. Secondly, for the case $\text{upp}(\psi) = 1$, we have $|\text{Res}(Q|_\psi)| \leq 1$. Let $s = \max\{i | l_i \neq h_i\}$, by Property 1, there exists at most one integer $p \in \mathbb{N}[l_s, h_s]$ such that

- $\text{upp}([l_1, h_1], \dots, [l_s, p-1], \dots, [l_n, h_n]) = 0$,
- $\text{upp}([l_1, h_1], \dots, [p, p], \dots, [l_n, h_n]) = 1$,
- $\text{upp}([l_1, h_1], \dots, [p+1, h_s], \dots, [l_n, h_n]) = 0$.

If $\text{upp}([l_1, h_1], \dots, [p, p], \dots, [l_n, h_n]) = 0$ follows for every $p \in \mathbb{N}[l_s, h_s]$, we have $\text{Res}(Q|_\psi) = \emptyset$. Otherwise, it is able to calculate the only integer $p \in \mathbb{N}[l_s, h_s]$ by a multi-head binary search (introduced in Section 4.4) in $O(\log |Q|)$ time. Then apply the same procedure to the next unfixed position of $[l_1, h_1], \dots, [p, p], \dots, [l_n, h_n]$, and continue this process iteratively on other positions as needed. If this process finally obtains $\psi^* = [[l_1^*, h_1^*], \dots, [l_n^*, h_n^*]]$ such that $\text{upp}(\psi^*) = 1$ and $\forall 1 \leq i \leq n, l_i^* = h_i^*$, then (l_1^*, \dots, l_n^*) is the only tuple in $\text{Res}(Q|_\psi)$. Otherwise, we have $\text{Res}(Q|_\psi) = \emptyset$. The time it takes to decide whether $\text{Res}(Q|_\psi) = \emptyset$ and calculate ψ^* (if exists) is at most $O(n \log |Q|)$, which is $O(\log |Q|)$ in data complexity. □

0.4 Proof of Lemma 8.

PROOF. For the convenience of the proof, a dummy element $+\infty$ is appended to the end of each sequence, that is, $A_i[A_i] = +\infty$ for each $1 \leq i \leq k$. If $F(|A_1|, \dots, |A_k|) \leq T$, the algorithm obviously returns p^* . Then for the case where $F(|A_1|, \dots, |A_k|) > T$, we show that Algorithm 4 maintains the invariant

$$A_i[l_i] \leq p^* \leq A_i[r_i], \quad \forall 1 \leq i \leq k, \quad (3)$$

and then prove that Algorithm 4 correctly returns p^* after at most $O(\log \sum_{i=1}^k |A_i|)$ iterations.

Base case. Initially, for each $1 \leq i \leq k$ we set $l_i = 0$ and $r_i = |A_i|$. Since each A_i is sorted,

$$A_i[0] \leq \min A_i \leq p^* \leq \max A_i \leq A_i[n], \quad (4)$$

then the invariant holds before the first iteration.

Inductive step. Suppose at the start of some iteration we have $A_i[l_i] \leq p^* \leq A_i[r_i]$ for any $1 \leq i \leq k$. We aim to demonstrate that the invariant is preserved after the update performed in the current iteration. Without loss of generality, we assume that for any $1 \leq i \leq k$, $r_i - l_i > 1$ satisfies. (If some i already has $r_i - l_i \leq 1$ and the algorithm has not yet terminated, then by the invariant we have $A_i[l_i] < p^* \leq A_i[r_i]$ and $N_i(A_i[r_i]) \leq N_i(A_i[l_i] + 1) = r_i$, then $N_i(p^*) = r_i = m_i$, and the problem reduces to the remaining $k-1$ sequences.) In each iteration, the algorithm modifies exactly one of the two bounds: it either updates the lower bound $l_{i_{\min}}$ for some index i_{\min} , or updates the upper bound $r_{i_{\max}}$ for some index i_{\max} . We discuss the two cases:

Case 1: $F(m_1, \dots, m_k) \leq T$.

Let $p = A_{i_{\min}}[m_{i_{\min}}] = \min_{r_i - l_i > 1} A_i[m_i]$. Since

$$N_{i_{\min}}(p) = |\{x \in A_{i_{\min}} \mid x < A_{i_{\min}}[m_{i_{\min}}]\}| \leq m_{i_{\min}}, \quad (5)$$

and for each $i \neq i_{\min}$,

$$N_i(p) \leq N_i(A_i[m_i]) = |\{x \in A_i \mid x < A_i[m_i]\}| \leq m_i, \quad (6)$$

it follows that

$$F(N_1(p), \dots, N_k(p)) \leq F(m_1, \dots, m_k) \leq T. \quad (7)$$

Hence $p^* \geq p$, then after updating $l_{i_{\min}} \leftarrow m_{i_{\min}}$, we still have $A_{i_{\min}}[l_{i_{\min}}] \leq p^* \leq A_{i_{\min}}[r_{i_{\min}}]$. For all $i \neq i_{\min}$, neither l_i nor r_i changes, so the invariant is preserved after the update.

Case 2: $F(m_1, \dots, m_k) > T$.

Let $p = A_{i_{\max}}[m_{i_{\max}}] = \max_{r_i - l_i > 1} A_i[m_i]$. Since

$$N_{i_{\max}}(p+1) = |\{x \in A_{i_{\max}} | x \leq A_{i_{\max}}[m_{i_{\max}}]\}| \geq m_{i_{\max}}, \quad (8)$$

and for each $i \neq i_{\max}$,

$$N_i(p+1) \geq N_i(A_i[m_i] + 1) = |\{x \in A_i | x \leq A_i[m_i]\}| \geq m_i, \quad (9)$$

it follows that

$$F(N_1(p+1), \dots, N_k(p+1)) \geq F(m_1, \dots, m_k) > T. \quad (10)$$

Thus $p^* < p+1$ which implies $p^* \leq p$, then after updating $r_{i_{\max}} \leftarrow m_{i_{\max}}$, we still have $A_{i_{\max}}[l_{i_{\max}}] \leq p^* \leq A_{i_{\max}}[r_{i_{\max}}]$. For all $i \neq i_{\max}$, neither l_i nor r_i changes, so the invariant is preserved after the update.

Convergence and termination. In each iteration, an update is performed on sequence A_i only if $r_i - l_i > 1$, and every such update reduces the length of the interval by at least one. Consequently, the length of each interval $[l_i, r_i]$ shrinks to at most 2 after at most $O(\log |A_i|)$ updates on A_i . Since the number of sequences is a constant, the total number of iterations of the algorithm is bounded by $O(\log \sum_{i=1}^k |A_i|)$.

Now we show that when Algorithm 4 terminates, it returns p^* correctly. If it terminates at Line 4 or Line 14, which implies

$$F(N_1(A_i[l_i] + 1), \dots, N_k(A_i[l_i] + 1)) > T, \quad (11)$$

it follows that $p^* \leq A_i[l_i]$, then by the invariant $p^* \geq A_i[l_i]$, we have $p^* = A_i[l_i]$.

If Algorithm 4 terminates at line 16, which implies $r_i - l_i \leq 1$ and $F(N_1(A_i[l_i] + 1), \dots, N_k(A_i[l_i] + 1)) \leq T$ satisfies for every $1 \leq i \leq k$. Let $p = \min_{1 \leq i \leq k} A_i[r_i]$ and $q = \max_{1 \leq i \leq k} A_i[l_i]$, then

$$\begin{aligned} F(N_1(p), \dots, N_k(p)) &\leq F(N_1(A_1[r_1]), \dots, N_k(A_k[r_k])) \\ &\leq F(N_1(A_1[l_1] + 1), \dots, N_k(A_k[l_k] + 1)) \\ &\leq F(N_1(q+1), \dots, N_k(q+1)) \leq T. \end{aligned} \quad (12)$$

Thus $p^* \geq p$, and by the invariant, $p^* \leq p$, we conclude $p^* = p = \min_{1 \leq i \leq k} A_i[r_i]$.

In summary, since i_{\min} , i_{\max} , and F can be computed in $O(1)$ time, and computing each N_i takes at most $O(\log |A_i|)$ time, Algorithm 4 returns p^* in $O(\log \sum_{i=1}^k |A_i|)$ time. \square

0.5 Proof of Lemma 9.

PROOF. Let $\psi = [[l_1, h_1], \dots, [l_n, h_n]]$, since $s = \min\{i | l_i \neq h_i\}$ is a split position of ψ , it follows that

- (1) $\forall 1 \leq i < s, l_i = h_i$,
- (2) $\forall s < i \leq n, l_i = \min_Q(x_i)$ and $h_i = \max_Q(x_i)$.

If ψ_{left} is not an empty range, then $l_p \leq \text{mid}_p - 1$, which implies that ψ_{left} is a prefix range filter with split position s . Similarly, if ψ_{left} is not an empty range, then $\text{mid}_p + 1 \leq h_p$, which implies that ψ_{right} is a prefix range filter with split position s . Finally, since $\text{mid}_p \leq \text{mid}_p$, ψ_{mid} is a prefix range filter with split position s . \square

0.6 Proof of Lemma 11.

PROOF. If $\text{att}(Q_s) = \text{att}(Q)$, then $\text{Res}(Q|_\psi) \subseteq \text{Res}(Q_s|_\psi)$, which implies $|\text{Res}(Q|_\psi)| \leq |\text{Res}(Q_s|_\psi)|$. If $\text{att}(Q_s) \subsetneq \text{att}(Q)$, then

$$\begin{aligned} |\text{Res}(Q|_\psi)| &= |\text{Res}(Q_s|_\psi) \bowtie \text{Res}(Q_r|_\psi)| \\ &\leq \sum_{t \in \text{Res}(Q_s|_\psi)} |\text{Res}((Q_r \bowtie t)|_\psi)| \\ &\leq |\text{Res}(Q_s|_\psi)| \max_{t \in \text{Res}(Q_s|_\psi)} |\text{Res}((Q_r \bowtie t)|_\psi)| \quad (13) \\ &\leq |\text{Res}(Q_s|_\psi)| \cdot |\text{Res}(Q_r^*|_\psi)| \\ &\leq |\text{Res}(Q_s|_\psi)| \cdot \text{AGM}_c(Q_r^*|_\psi). \end{aligned}$$

\square

0.7 Queries in the Experiments.

Query 1: Q_A on TPC-DS+

```
SELECT *
FROM follow f, web_sales w1, web_sales w2, web_sales w3,
     web_sales w4
WHERE f.src = w1.ws_bill_customer_sk
      AND f.src = w2.ws_bill_customer_sk
      AND f.dst = w3.ws_bill_customer_sk
      AND f.dst = w4.ws_bill_customer_sk
      AND w1.ws_item_sk = w3.ws_item_sk
      AND w2.ws_item_sk = w4.ws_item_sk
```

Query 2: Q_Δ on Twitter

```
SELECT *
FROM follow A, follow B, follow C
WHERE A.dst = B.src AND B.dst = C.src AND C.dst = A.src
```

Query 3: Q_\S on Twitter

```
SELECT *
FROM follow A, follow B, follow C, follow D, follow E,
     follow F
WHERE A.dst = B.src AND B.dst = C.src
      AND C.dst = D.src AND D.dst = A.src
      AND E.src = A.src AND E.dst = B.dst
      AND F.dst = A.dst AND F.src = D.src
```